

from zero to query

a sql primer

oskar 2023-09-13

a note on `sqlite`

- small (<2mb)
- open source
- serverless
- self-contained
- fast
- complete
- in-memory
- cross-platform
- ubiquitous



sqlite commands



- these are not **sql** commands!
- they start with a **'.'**
- they operate on the environment, not the data
- examples:
 - **.quit**
 - **.open <path-to-database>**
 - **.show**
 - **.help**
 - **.cd <directory>**
 - **.shell CMD ARGS...**

```
.open data/sqlite-sakila.db  
.tables
```

```
.header ON  
.mode qbox
```

.tables

```
sqlite> .tables
```

actor

address

category

city

country

customer

customer_list

```
sqlite> █
```

film

film_actor

film_category

film_list

film_text

inventory

language

payment

rental

sales_by_film_category

sales_by_store

staff

staff_list

store

sql - a fundamental data tool

- database management
- data pipeline engineering
- data modeling
- data designing
- **big** data (parallel, distributed)
- data querying
- data analytics

data definition	data management	data querying	data control	transaction control
to operate on entire tables	to operate on table values, rows, columns	to fetch data from tables	to control access to schemas + tables	for transactional atomicity, dev
CREATE	INSERT	SELECT	GRANT	COMMIT
DROP	UPDATE		REVOKE	ROLLBACK
ALTER	DELETE			SAVE POINT
TRUNCATE				

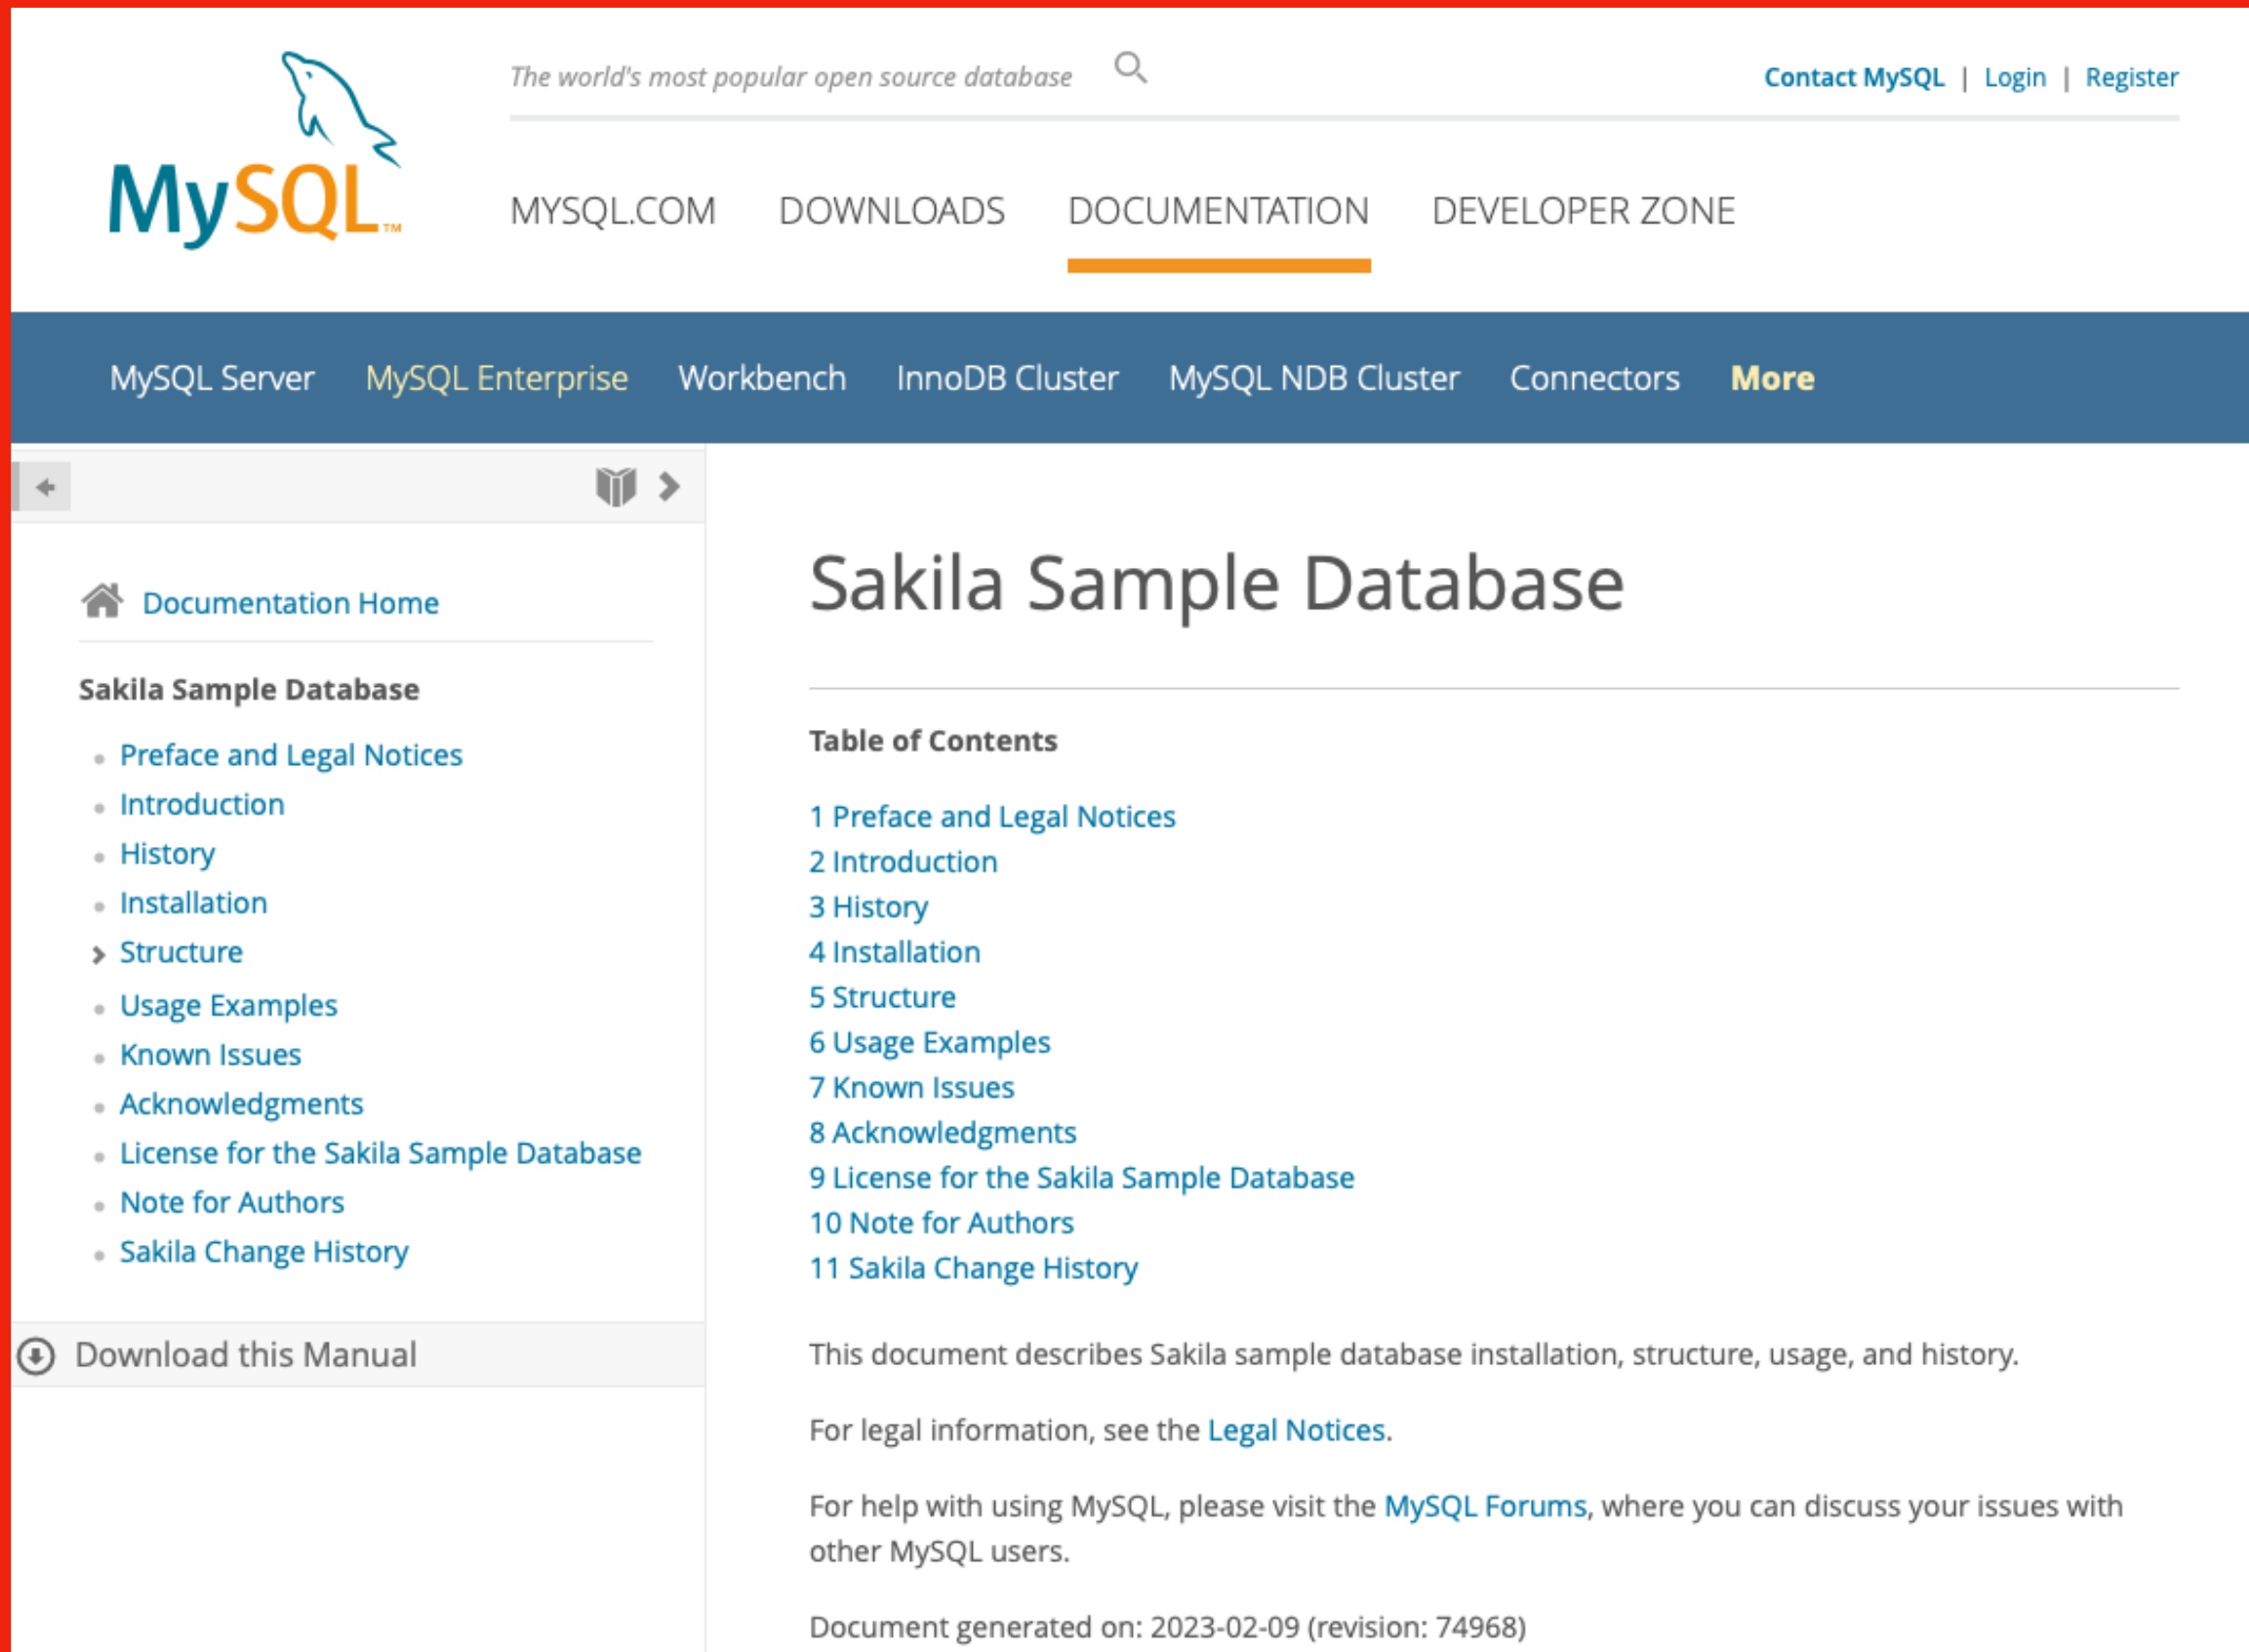
data definition	data management	data querying	data control	transaction control
to operate on entire tables	to operate on table cells, rows, columns	to fetch data from tables	to control access to schemas + tables	for transactional atomicity, dev
CREATE	INSERT	SELECT	GRANT	COMMIT
DROP	UPDATE		REVOKE	ROLLBACK
ALTER	DELETE			SAVE POINT
TRUNCATE				

sql commands

- run on the database
- operate on the data tables
- output a table
- start with **SELECT** ... clause
- contain a **FROM** ... clause
- end with a **;**
- you can add comments with **-- a comment**
- example:
SELECT name FROM category; -- film categories

the sakila training data

- classic, fictional data
- dvd rental company
- 20 relational tables:
 - normalised, no repetition
- stores
- inventory
- films
- film casting
- actors
- film ratings



The screenshot shows the MySQL website's documentation page for the Sakila Sample Database. The page is titled "Sakila Sample Database" and features a "Table of Contents" with 11 items: 1 Preface and Legal Notices, 2 Introduction, 3 History, 4 Installation, 5 Structure, 6 Usage Examples, 7 Known Issues, 8 Acknowledgments, 9 License for the Sakila Sample Database, 10 Note for Authors, and 11 Sakila Change History. The page also includes a "Download this Manual" button and a footer stating "Document generated on: 2023-02-09 (revision: 74968)".

MySQL
The world's most popular open source database

MySQL.COM DOWNLOADS DOCUMENTATION DEVELOPER ZONE

MySQL Server MySQL Enterprise Workbench InnoDB Cluster MySQL NDB Cluster Connectors More

Sakila Sample Database

Table of Contents

- 1 Preface and Legal Notices
- 2 Introduction
- 3 History
- 4 Installation
- 5 Structure
- 6 Usage Examples
- 7 Known Issues
- 8 Acknowledgments
- 9 License for the Sakila Sample Database
- 10 Note for Authors
- 11 Sakila Change History

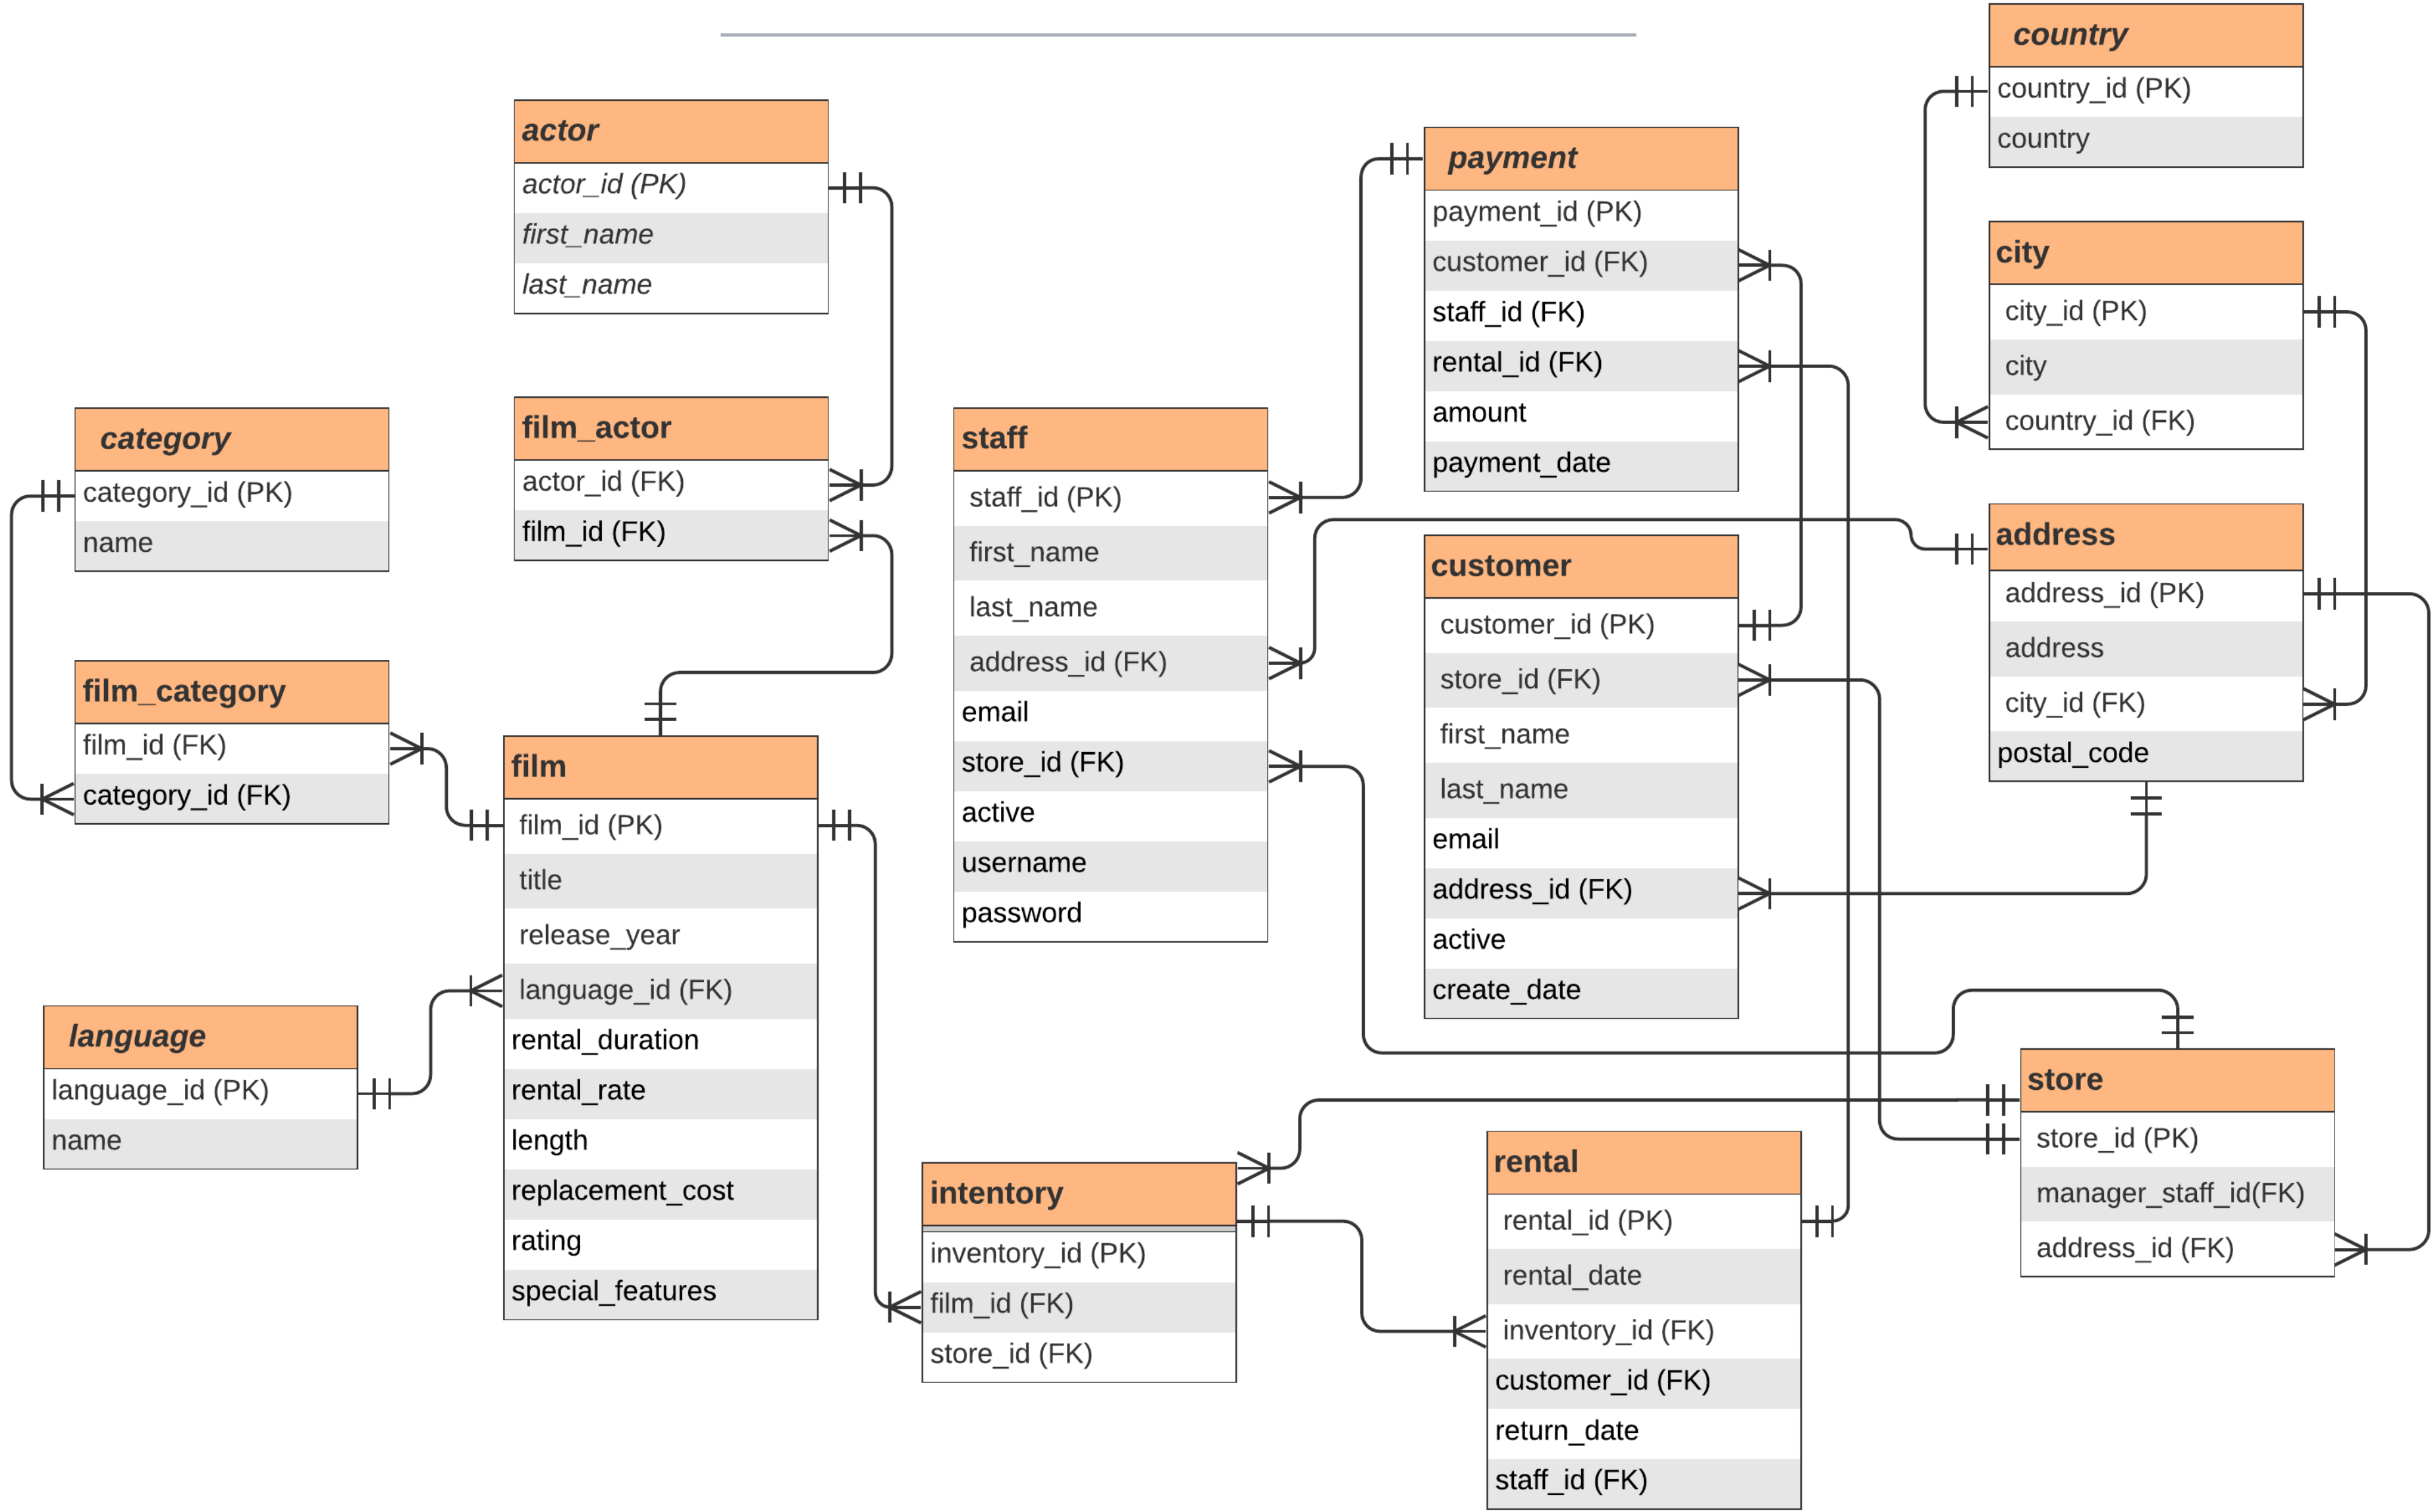
This document describes Sakila sample database installation, structure, usage, and history.

For legal information, see the [Legal Notices](#).

For help with using MySQL, please visit the [MySQL Forums](#), where you can discuss your issues with other MySQL users.

Document generated on: 2023-02-09 (revision: 74968)

SQLite3 Sakila Sample Database ERD



today's objective:

“which **top 10** *actors* were *rented* out
the greatest number of times, counting
only ‘**R**’ *rated films* made in **2006?**”

today's plan:

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

```
- SELECT {columns} FROM {table};  
  INNER JOIN {table_2} ON {col1}={col2}  
  WHERE {condition}  
  GROUP BY {columns}  
  HAVING {condition}  
  ORDER BY {columns}  
  LIMIT {num}
```

what do the tables contain?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

- `SELECT {columns} FROM {table};`
 `INNER JOIN {table_2} ON {col1}={col2}`
 `WHERE {condition}`
 `GROUP BY {columns}`
 `HAVING {condition}`
 `ORDER BY {columns}`
 `LIMIT {num}`

SELECT {columns} FROM {table};

- `SELECT * FROM staff;`
-- returns all columns and all rows from the staff table
- `SELECT title, rating FROM film;`
-- returns title and rating (in order) from the table film
- `SELECT c.first_name AS customer_name FROM customer c;`
-- sets an alias for table customer, renames column to 'customer_name'
- `SELECT title, replacement_cost/rental_rate AS break_even_count FROM film;`
-- returns the number of rentals a film needs to break even
- `SELECT DISTINCT a.last_name FROM actor a;`
-- returns all the first names in the actor table, with no duplicates

exercises

- show all the columns of the `category` table
- rename the `name` column to `category_name` in the output

SELECT {aggregate function} FROM {table};

- SELECT COUNT(*) AS num_records FROM actor;
-- returns the number of rows in table actor, names the output 'num_records'
- SELECT COUNT(DISTINCT rating) FROM film;
-- returns a count of distinct values in the rating column
- SELECT AVG(replacement_cost) AS avg_cost FROM film;
-- returns the average replacement cost of a film
- SELECT AVG(rental_rate + replacement_cost) AS average_total_cost FROM film;
-- returns the average rate of rental from film table
- SELECT MAX(rental_rate) AS highest_rental_rate FROM film;
-- returns the most expensive rental_rate from film
- SELECT MIN(length) AS shortest_length FROM film;
-- returns the length of the shortest film

exercises

- what is the maximum number of times that a film needs to be rented out to break even?
- what is the average number of times that a film needs to be rented out to break even?

that's too many rows!

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

```
- SELECT {columns} FROM {table}
  INNER JOIN {table_2} ON {col1}={col2}
  WHERE {condition}
  GROUP BY {columns}
  HAVING {condition}
  ORDER BY {columns}
  LIMIT {num}
;
```

SELECT ... FROM ... LIMIT ...;

- `SELECT * FROM film LIMIT 5;`
-- returns all columns of 5 unspecified rows from film
- `SELECT * FROM category LIMIT 5;`
-- returns 5 unspecified rows of all columns from category
- `SELECT title, release_year FROM film LIMIT 15;`
-- returns 15 unspecified rows of two columns from film table
- `SELECT r.rental_id, r.rental_date FROM rental r LIMIT 10;`
-- returns rental id and date of rental for 10 unspecified rows
- `SELECT first_name||' '||last_name AS fullname FROM actor LIMIT 10;`
-- returns the full names of 10 unspecified actor

exercises

- show the first 5 records of any table, using a **LIMIT** clause
- output the full name of 5 **customers**, using **||** to paste strings

but i only want the most extreme rows!

“which top 10 actors were rented out
the greatest number of times, counting
only ‘R’ rated films made in 2006?”

```
- SELECT {columns} FROM {table}  
  INNER JOIN {table_2} ON {col1}={col2}  
  WHERE {condition}  
  GROUP BY {columns}  
  HAVING {condition}  
  ORDER BY {columns}  
  LIMIT {num}  
;
```

SELECT ... FROM ... ORDER BY ... LIMIT ...;

- SELECT * FROM payment ORDER BY payment_date LIMIT 7;
-- return the earliest 7 payments in the payment table
- SELECT * FROM payment ORDER BY payment_date DESC LIMIT 7;
-- return the latest 7 payments in the payment table
- SELECT * FROM payment ORDER BY amount DESC LIMIT 5;;
-- return only the top 5 highest payment amounts from the payment table

but i only want specific rows!

“which top 10 actors were rented out
the greatest number of times, counting
only ‘R’ rated films made in 2006?”

```
- SELECT {columns} FROM {table}  
  INNER JOIN {table_2} ON {col1}={col2}  
  WHERE {condition}  
  GROUP BY {columns}  
  HAVING {condition}  
  ORDER BY {columns}  
  LIMIT {num}  
;
```


comparison operators

comparison syntax	meaning
{column} = {expression}	column value is equal to expression value
{column} <> {expression}	column value is not equal to expression value
{column} != {expression}	column value is not equal to expression value
{column} < {expression}	column value is less than expression value
{column} <= {expression}	column value is less than or equal to expression value
{column} > {expression}	column value is greater than expression value
{column} >= {expression}	column value is greater than or equal to expression value
{column} IN ({exp1,exp2,...})	column value is one of exp1, exp2, ...
{column} LIKE '%expr%'	(string) column contains substring 'expr'
{column} BETWEEN {exp1} AND {exp2}	{exp1} <= column value <= {exp2}

SELECT {column} FROM {table} WHERE {cond}

- SELECT * FROM actor WHERE LENGTH(last_name) = 3;
-- returns only records of actors whose last name is three characters
- SELECT title AS name FROM film f WHERE rating <>'R' AND;
-- returns only rows where the value in column1 is not {expression}
- SELECT title AS film_name, rental_rate FROM film WHERE rental_rate<=1.0;
-- returns titles of films whose rental price is at most £1
- SELECT first_name FROM staff WHERE store_id=2;
-- returns the first names of staff at store with id 2

exercises

- show the `first_names` of inactive customers (active = '0')
- how many `payments` have `amounts` greater than \$10.00?

comparison operators

comparison syntax	meaning
{column} = {expression}	column value is equal to expression value
{column} <> {expression}	column value is not equal to expression value
{column} != {expression}	column value is not equal to expression value
{column} < {expression}	column value is less than expression value
{column} <= {expression}	column value is less than or equal to expression value
{column} > {expression}	column value is greater than expression value
{column} >= {expression}	column value is greater than or equal to expression value
{column} IN ({exp1,exp2,...})	column value is one of exp1, exp2, ...
{column} LIKE '%expr%'	(string) column contains substring 'expr'
{column} BETWEEN {exp1} AND {exp2}	{exp1} <= column value <= {exp2}

SELECT ... FROM ... WHERE ...;

- `SELECT * FROM rental WHERE rental_date BETWEEN '2005-08-16' AND '2005-08-17';`
-- returns only rentals occurring in
- `SELECT * FROM payment WHERE amount IN (7.98, 8.97);`
-- returns info on all payments of a specific amount
- `SELECT * FROM city WHERE city LIKE '0k%';`
-- returns info on all cities whose name begins with '0k'
- `SELECT last_name AS full_name FROM customer WHERE first_name LIKE 'AL%';`
-- returns the last name of all customers whose first name begins with 'AL'

exercises

- what `city` names begin with a 'Q'?
- which `actors'` `first_names` end with 'K'?

how can i aggregate groups of rows into a single row?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

```
– SELECT {columns} FROM {table};  
  INNER JOIN {table_2} ON {col1}={col2}  
  WHERE {a_condition}  
  GROUP BY {columns}  
  HAVING {a_condition}  
  ORDER BY {columns}  
  LIMIT {num}
```

SELECT {col} FROM {tab} GROUP BY {col};

- SELECT city_id, COUNT(*) AS num_address FROM address GROUP BY city_id;
-- return number of addresses in each city in address table
- SELECT rating, AVG(length) AS avg_len FROM film GROUP BY rating ORDER BY avg_len;
-- returns the average length of a movie in each rating category
- SELECT country_id, COUNT(*) AS num_cities
FROM city
GROUP BY country_id
ORDER BY num_cities DESC
LIMIT 5;
-- return top 5 country ids, by number of cities assigned to each

how do i report only some aggregated groups?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

```
- SELECT {columns} FROM {table};  
  INNER JOIN {table_2} ON {col1}={col2}  
  WHERE {condition}  
  GROUP BY {column}  
  HAVING {condition}  
  ORDER BY {columns}  
  LIMIT {num}  
  ;
```


SELECT ... FROM ... GROUP BY ... HAVING ...;

- `SELECT col1, COUNT(*) AS num FROM table GROUP BY col1 HAVING num>9;`
-- count instances of each value of col1, but only output rows with count>9
- `SELECT rating, AVG(length) AS len FROM film GROUP BY rating HAVING len<115;`
-- the film rating categories with average length of film under 115 minutes
- `SELECT actor_id, COUNT(*) AS n FROM film_actor GROUP BY actor_id HAVING n<15;`
-- which actor ids have appeared in fewer than 15 films?

but my information is spread over two tables!

“which top 10 actors were rented out
the greatest number of times, counting
only ‘R’ rated films made in 2006?”

```
- SELECT {columns} FROM {table};  
  INNER JOIN {table_2} ON {col1}={col2}  
  WHERE {condition}  
  GROUP BY {columns}  
  HAVING {condition}  
  ORDER BY {columns}  
  LIMIT {num}  
;
```

JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

...

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

...

JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

...

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

...

JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

...

Turkey

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

...

JOIN

city

city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

Turkey

Yemen

...

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

...

we want this

city-and-country

city_id	city	country
1	?	?
2	?	?
3	?	?
4	?	?
5	?	?
6	?	?
7	?	?
8	?	?

...

we want this

city-and-country

city_id	city	country
1	A Corua (La Corua)	Spain
2	Abha	Saudi Arabia
3	Abu Dhabi	United Arab Emirates
4	Acua	Mexico
5	Adana	Turkey
6	Addis Abeba	Ethiopia
7	Aden	Yemen
8	Adoni	India

...

so we add a JOIN to the WHERE clause

city


city_id	city	country_id	last_update
1	A Corua (La Corua)	87	2021-03-06 15:51:49
2	Abha	82	2021-03-06 15:51:49
3	Abu Dhabi	101	2021-03-06 15:51:49
4	Acua	60	2021-03-06 15:51:49
5	Adana	97	2021-03-06 15:51:49
6	Addis Abeba	31	2021-03-06 15:51:49
7	Aden	107	2021-03-06 15:51:49
8	Adoni	44	2021-03-06 15:51:49

...

country

country_id	country	last_update
1	Afghanistan	2021-03-06 15:51:49
2	Algeria	2021-03-06 15:51:49
3	American Samoa	2021-03-06 15:51:49
4	Angola	2021-03-06 15:51:49
5	Anguilla	2021-03-06 15:51:49
6	Argentina	2021-03-06 15:51:49
7	Armenia	2021-03-06 15:51:49
8	Australia	2021-03-06 15:51:49
9	Austria	2021-03-06 15:51:49

...



```
SELECT
    city_id, city.city, country.country
FROM
    city
    INNER JOIN country ON city.country_id=country.country_id
;
```

so we add a JOIN to the WHERE clause

city-and-country

=

city_id	city	country
1	A Corua (La Corua)	Spain
2	Abha	Saudi Arabia
3	Abu Dhabi	United Arab Emirates
4	Acua	Mexico
5	Adana	Turkey
6	Addis Abeba	Ethiopia
7	Aden	Yemen
8	Adoni	India

...

```
SELECT
    city_id, city.city, country.country
FROM
    city
    INNER JOIN country ON city.country_id=country.country_id
;
```

```
SELECT ... FROM a INNER JOIN b ON a.key=b.key;
```

- SELECT a.city, b.country
FROM city a
INNER JOIN country b ON a.country_id=b.country_id
LIMIT 10
; -- output a table with city-country names
- SELECT f.title, f.length, l.name
FROM film f
INNER JOIN language l ON f.language_id=l.language_id
WHERE f.rating='R'
LIMIT 10; -- output a sample of films and the name of the language it is in


```
SELECT ... FROM a INNER JOIN b ON a.key=b.key;
```

```
SELECT
    f.title AS film_title,
    c.name AS category
FROM film f
    INNER JOIN film_category fc ON f.film_id=fc.film_id
    INNER JOIN category c ON fc.category_id=c.category_id
WHERE f.rating IN ('G', 'PG') AND f.length BETWEEN 85 AND 90
;
```

exercises

- what countries are the cities whose name starts with 'Q' in?

FROM a INNER JOIN b



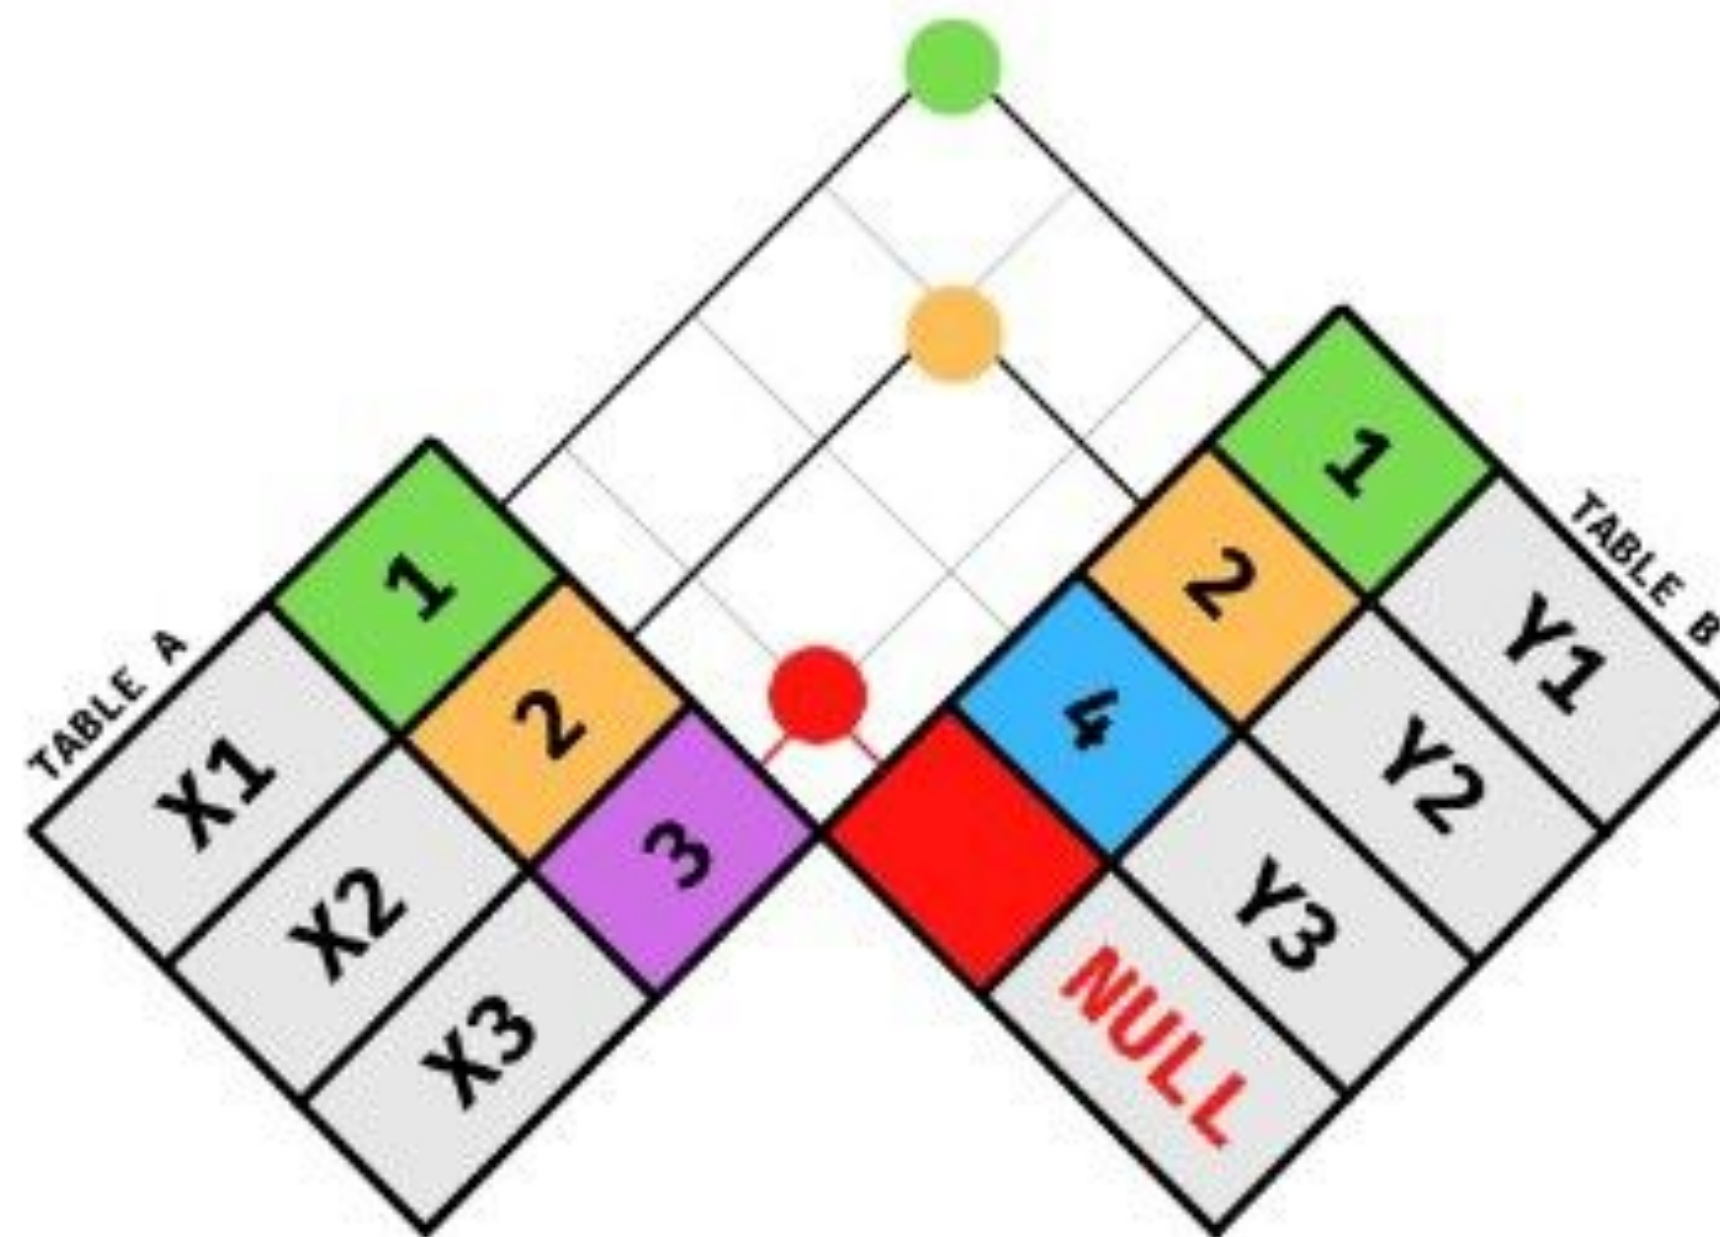
INNER JOIN



```
SELECT  
  <SELECT LIST>  
FROM    TABLE_A A  
INNER JOIN TABLE_B B  
  ON A.KEY = B.KEY
```

KEY	VAL_X	VAL_Y
1	X1	Y1
2	X2	Y2

LEFT OUTER JOIN



LEFT JOIN



```
SELECT  
    <SELECT LIST>  
FROM    TABLE_A A  
LEFT JOIN TABLE_B B  
ON A.KEY = B.KEY
```

KEY	VAL_X	VAL_Y
1	X1	Y1
2	X2	Y2
3	X3	NULL



@DATAWITHDANNY

Table 1

A
B
C

Table 2

A
B
D

INNER JOIN: show all matching records in both tables.

A	A
B	B

LEFT JOIN: show all records from left table, and any matching records from right table.

A	A
B	B
C	

RIGHT JOIN: show all records from right table, and any matching records from left table.

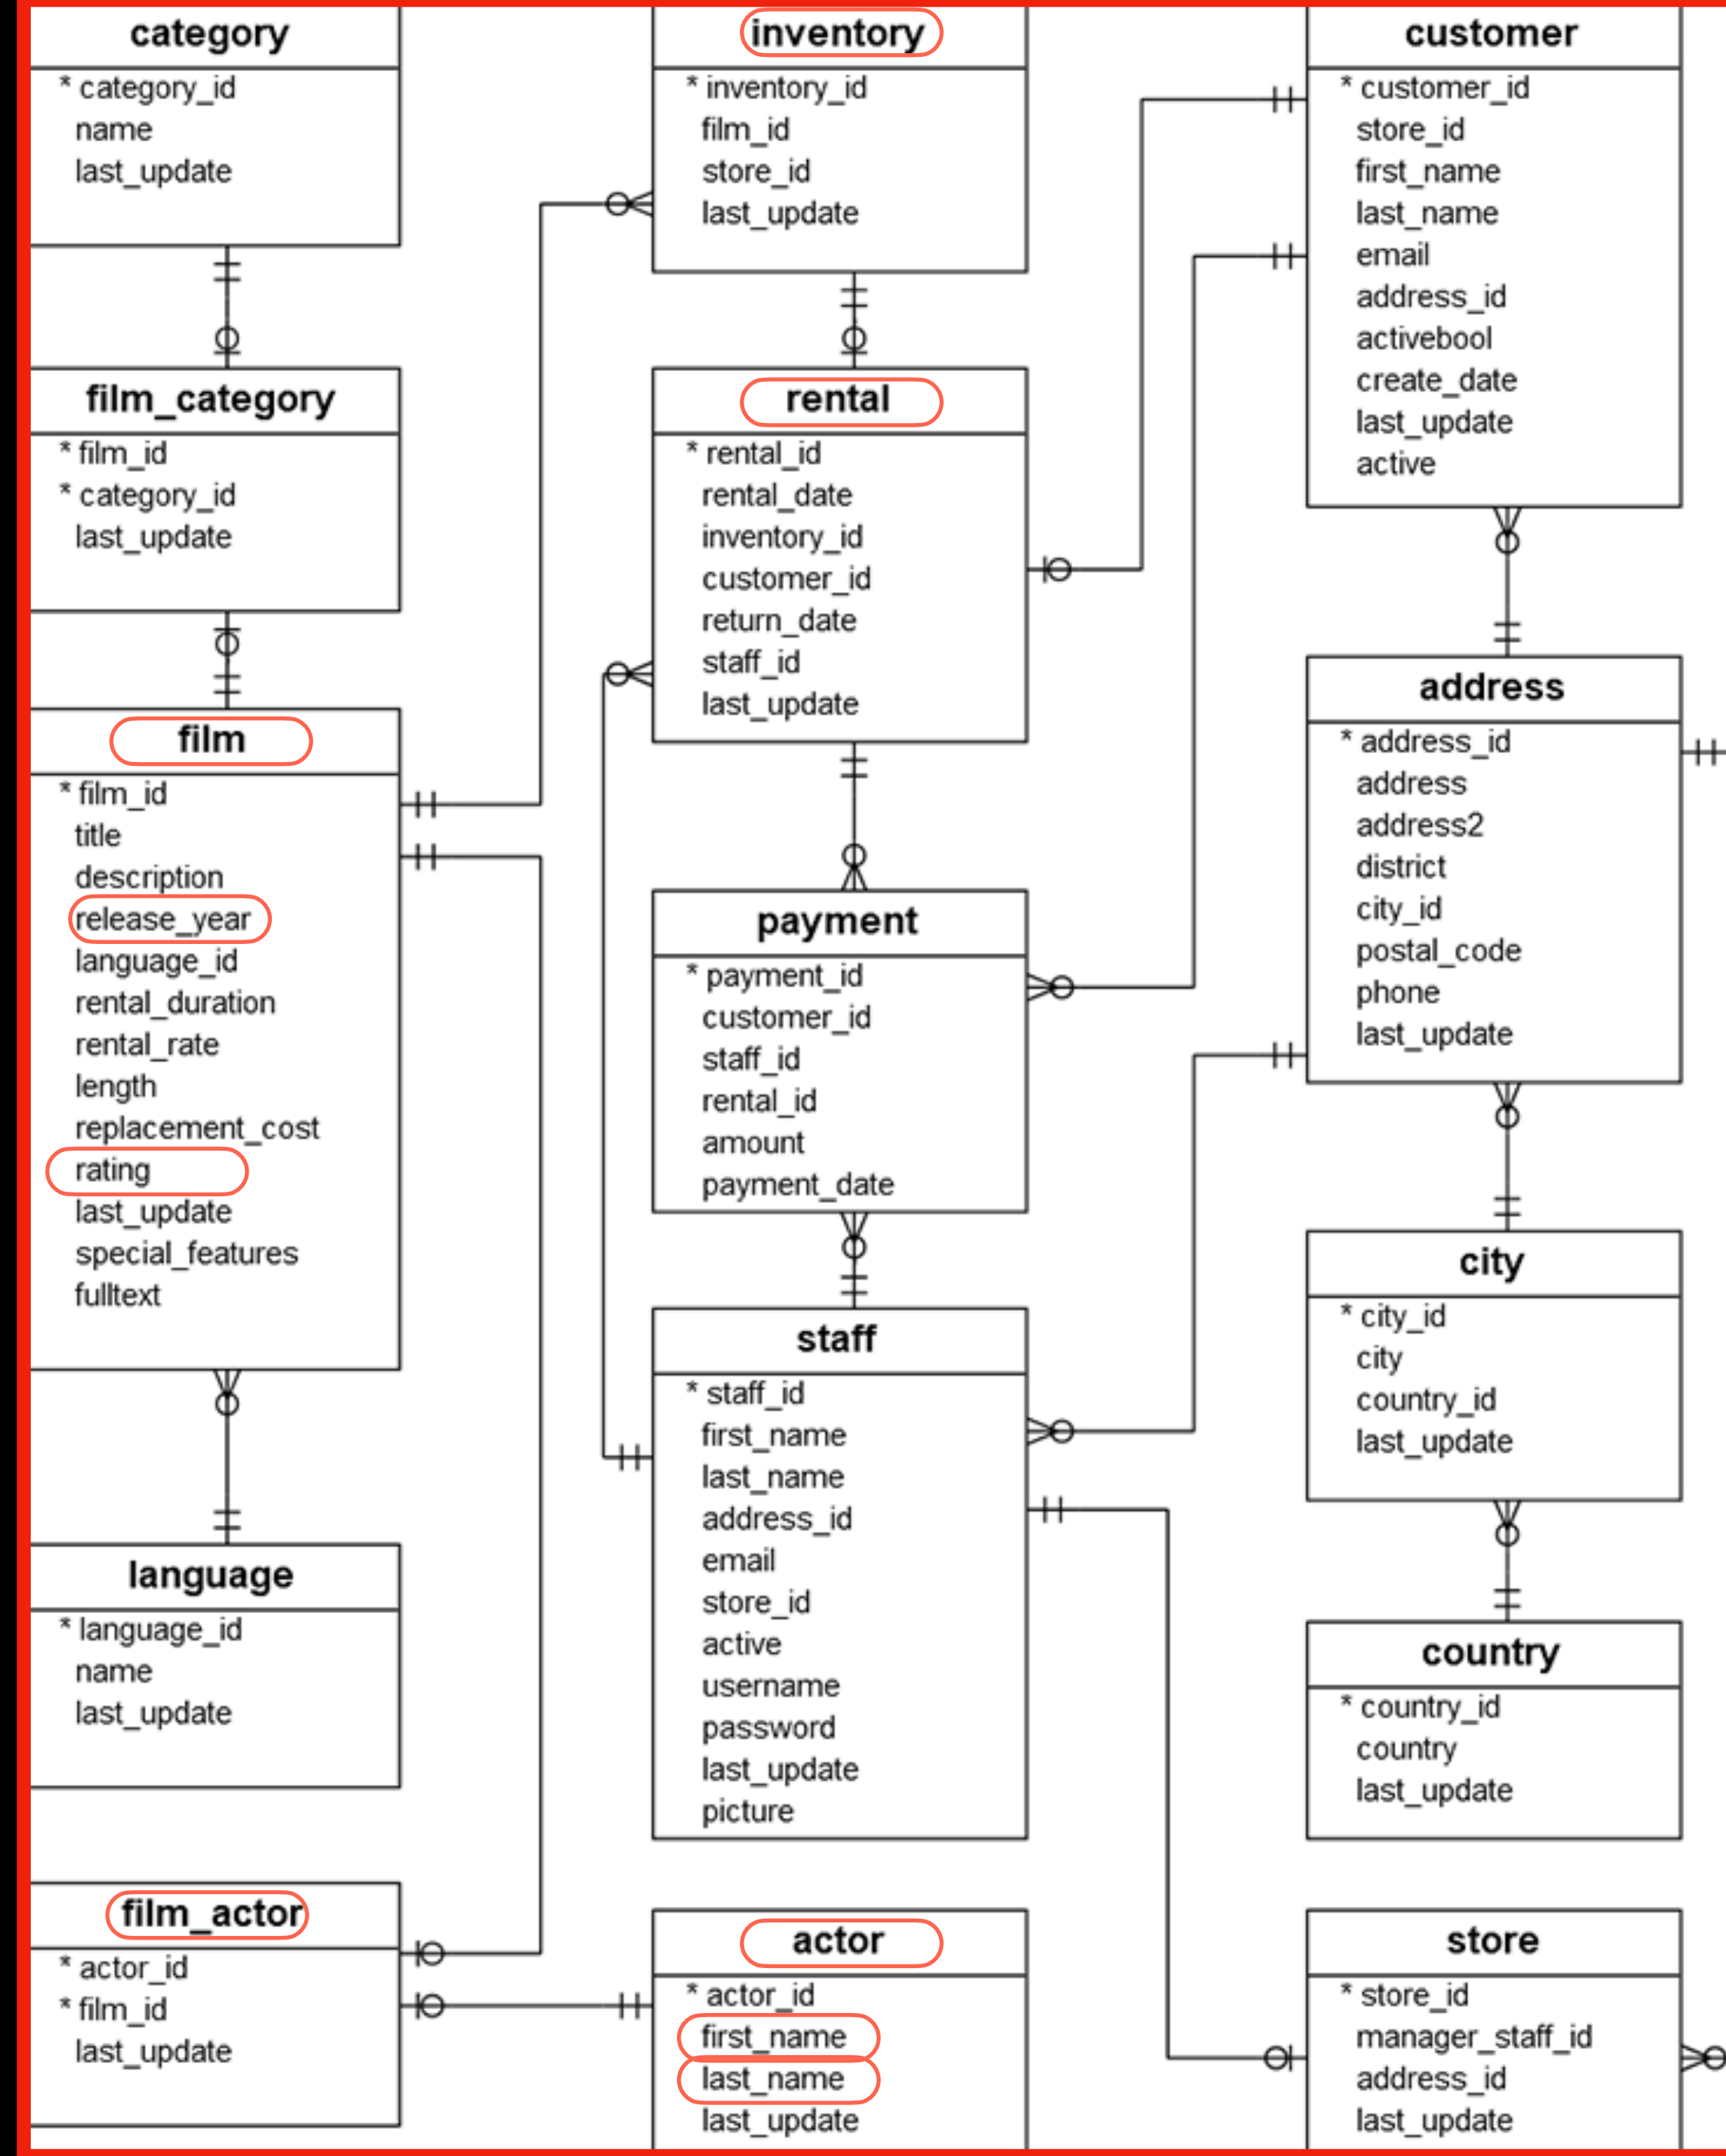
A	A
B	B
	D

FULL JOIN: show all records from both tables, whether there is a match or not.

A	A
B	B
C	
	D

today's objective:

“which **top 10** actors were rented out the greatest number of times, counting only ‘**R**’ rated films made in **2006**?”



how do i combine the components of a SELECT?

“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

```
– SELECT {columns} FROM {table};  
  INNER JOIN {table_2} ON {col1}={col2}  
  WHERE {condition}  
  GROUP BY {columns}  
  HAVING {condition}  
  ORDER BY {columns}  
  LIMIT {num}  
  ;
```

how the query
is written

SELECT ...

FROM + JOIN

WHERE ...

GROUP BY ...

HAVING ...

ORDER BY ...

LIMIT ...

how you should
think about it

FROM + JOIN

↓
WHERE

↓
GROUP BY

↓
HAVING

↓
SELECT

↓
ORDER BY

↓
LIMIT









your turn! compose a query to answer:

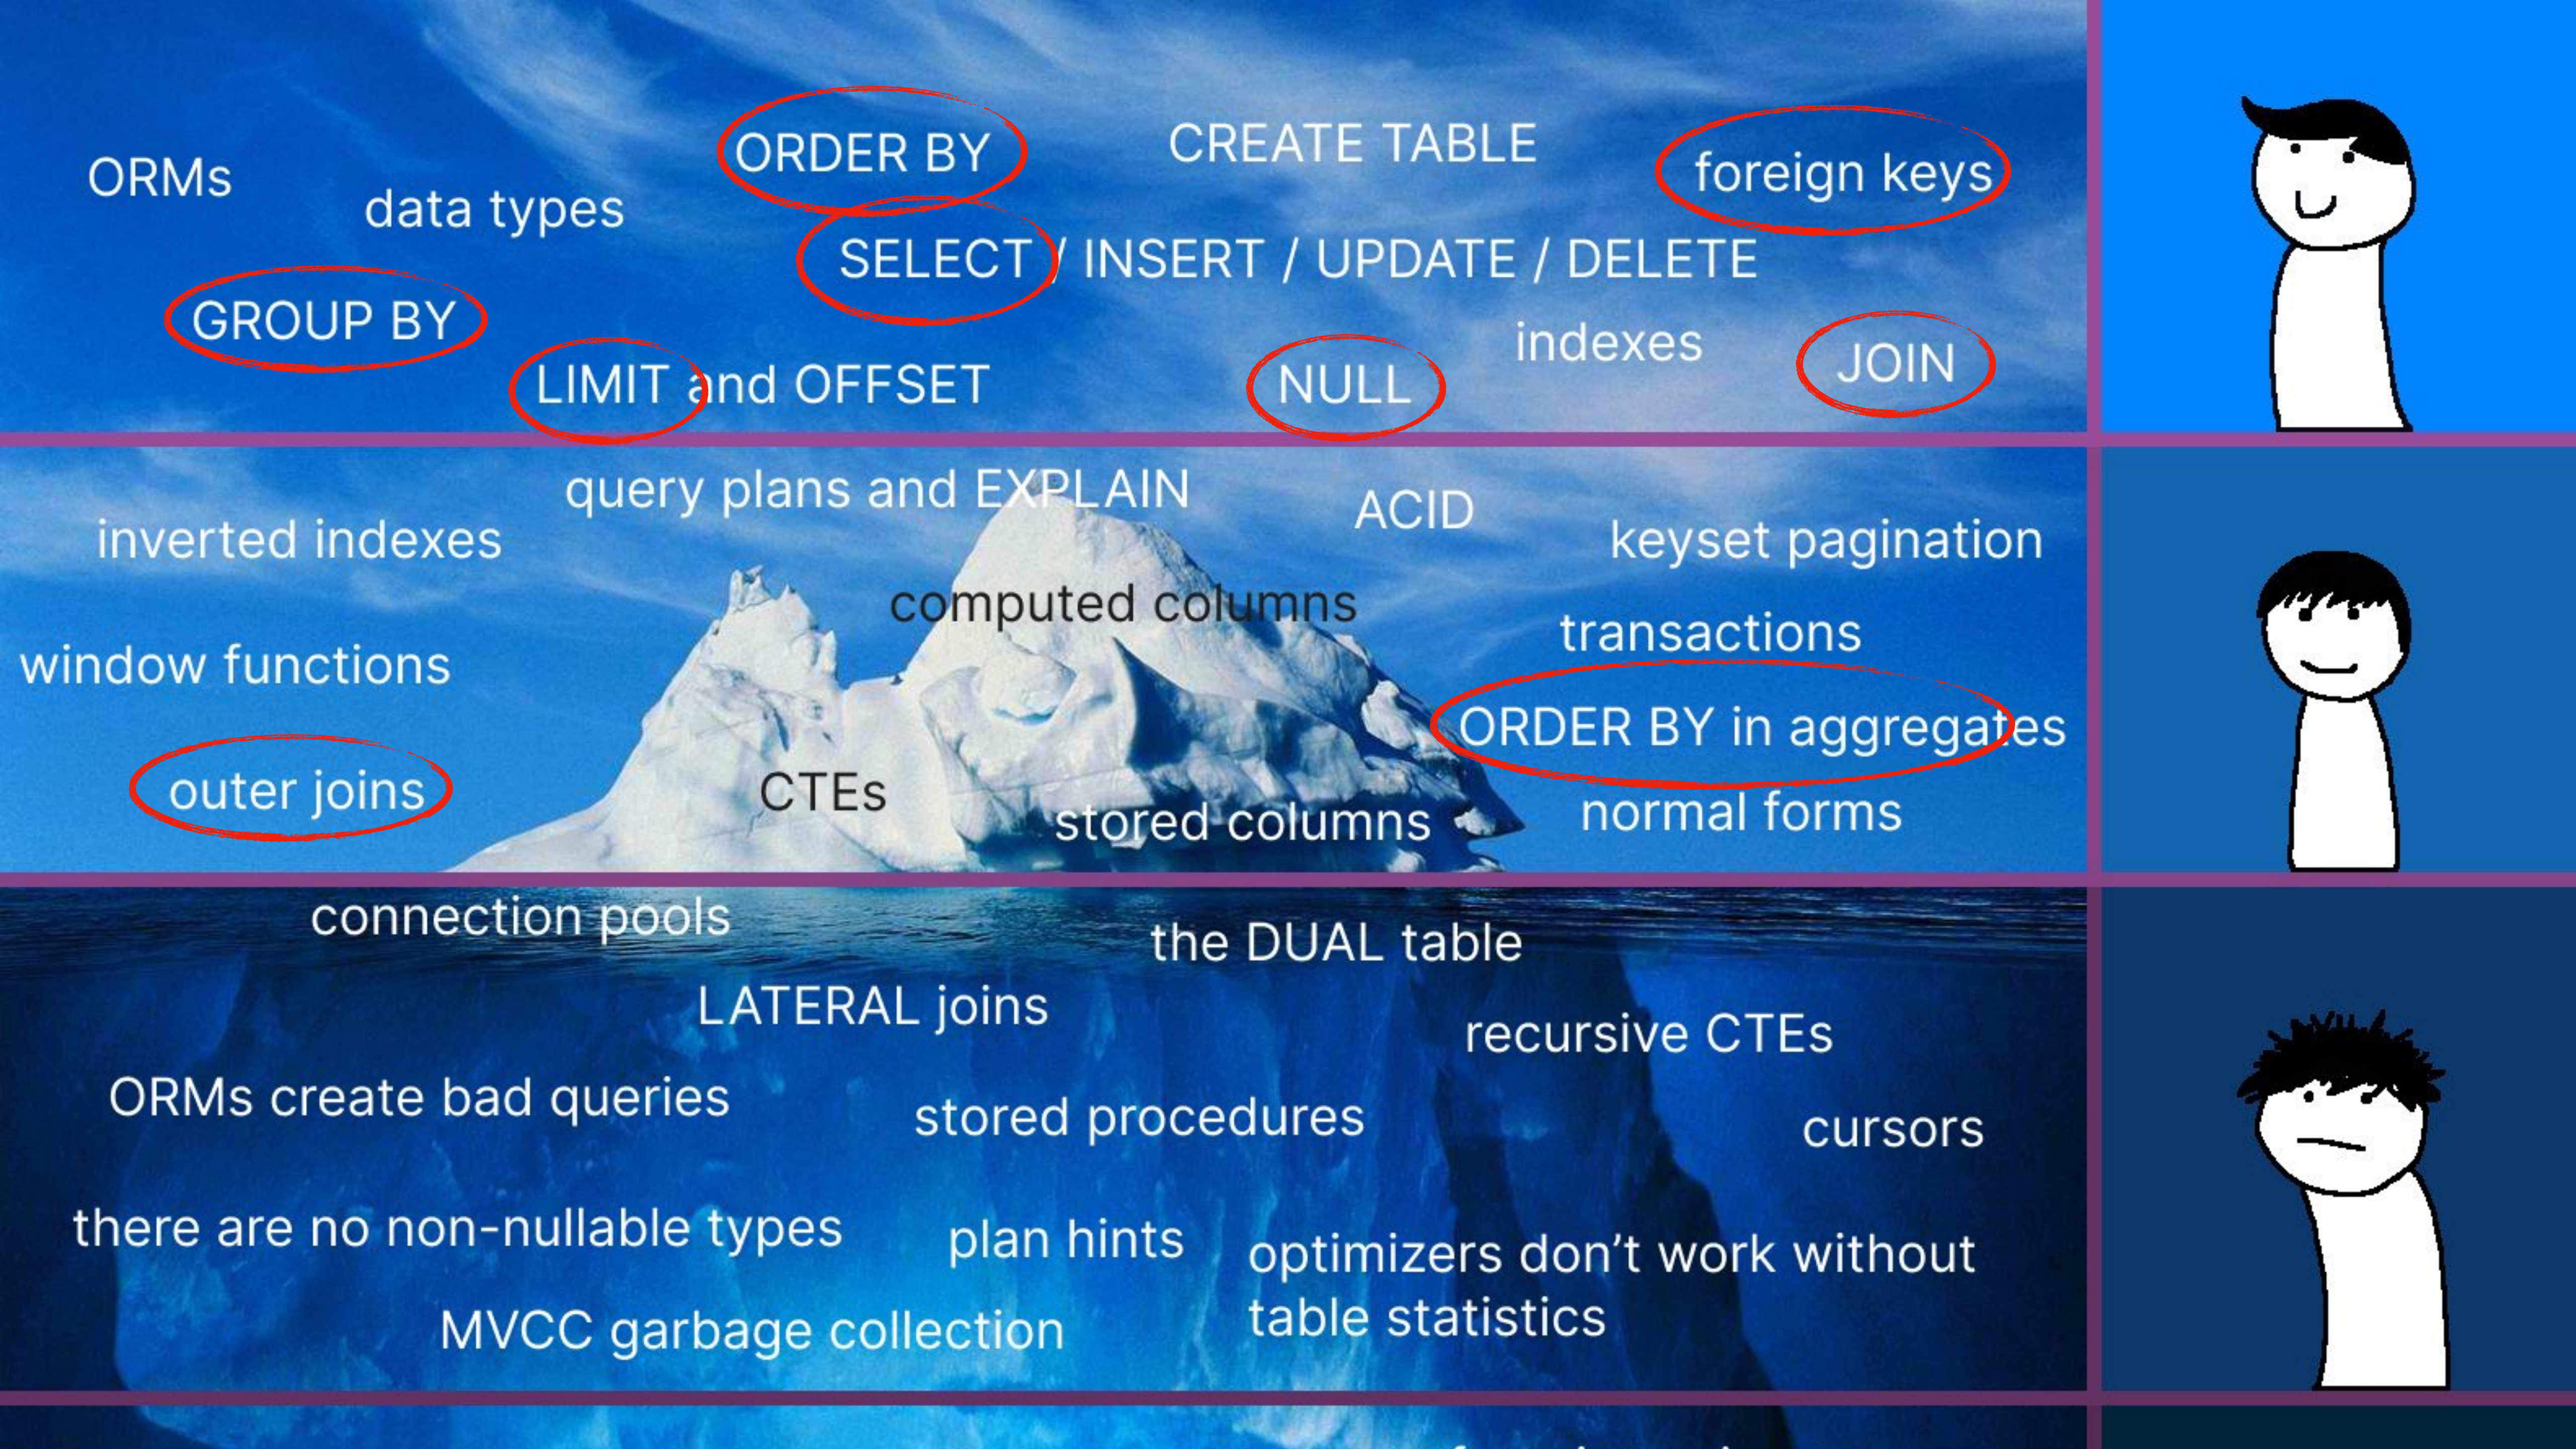
“which top 10 actors were rented out the greatest number of times, counting only ‘R’ rated films made in 2006?”

hint: structure of the solution

```
SELECT
    {}          AS actor_name,
    COUNT({}) AS num_rentals
FROM {}table1
    INNER JOIN {}table2 ON {}join-condition
    INNER JOIN {}table3 ON {}join-condition
    INNER JOIN {}table4 ON {}join-condition
    INNER JOIN {}table5 ON {}join-condition
WHERE {}row condition1
    AND {}row condition2
GROUP BY {}column1
ORDER BY {}column DESC
LIMIT {}num
;
```




ORMs	data types	ORDER BY	CREATE TABLE	foreign keys	
	GROUP BY	SELECT / INSERT / UPDATE / DELETE	indexes	JOIN	
	LIMIT and OFFSET	NULL			
inverted indexes	query plans and EXPLAIN	ACID	keyset pagination		
computed columns	transactions				
window functions	ORDER BY in aggregates				
outer joins	CTEs	stored columns	normal forms		
connection pools	the DUAL table				
LATERAL joins	recursive CTEs				
ORMs create bad queries	stored procedures	cursors			
there are no non-nullable types	plan hints	optimizers don't work without table statistics			
MVCC garbage collection					
COUNT(*) vs COUNT(1)	isolation levels	generator functions zip when cross joined	sharding		
serializable restarts require retry loops on all statements	zigzag join	phantom reads	triggers	MERGE	
grouping sets, cube, rollup	write skew	partial indexes			
denormalization	SELECT FOR UPDATE	NULLs in CHECK constraints are truthy			
transaction contention	sargability	timestamptz doesn't store a timezone	star schemas		
ascending key problem	ambiguous network errors	utf8mb4			
cost models don't reflect reality	'null':jsonb IS NULL = false	TPCC requires wait times			
DEFERRABLE INITIALLY IMMEDIATE		causal reverse			
EXPLAIN approximates SELECT COUNT(*)	MATCH PARTIAL foreign keys				
vectorized doesn't mean SIMD	NULLs are equal in DISTINCT but inequal in UNIQUE	volcano model			
join ordering is NP hard	database cracking	WCOJ			
learned indexes	XTID exhaustion				
the halloween problem	dee and dum	SERIAL is non-transactional			
fsyncgate	allballs	NULL	every sql operator is actually a join		



ORMs

data types

ORDER BY

CREATE TABLE

foreign keys

GROUP BY

SELECT / INSERT / UPDATE / DELETE

LIMIT and OFFSET

NULL

indexes

JOIN

inverted indexes

query plans and EXPLAIN

ACID

keyset pagination

computed columns

window functions

transactions

outer joins

ORDER BY in aggregates

CTEs

stored columns

normal forms

connection pools

the DUAL table

LATERAL joins

recursive CTEs

ORMs create bad queries

stored procedures

cursors

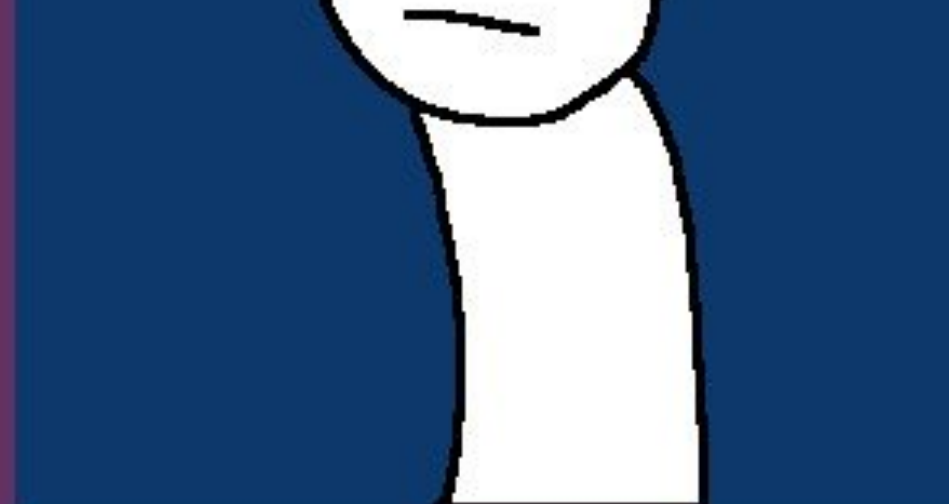
there are no non-nullable types

plan hints

optimizers don't work without
table statistics

MVCC garbage collection

there are no non-nullable types
plan hints
optimizers don't work without table statistics
MVCC garbage collection



COUNT(*) vs COUNT(1)
isolation levels
generator functions zip when cross joined
sharding
serializable restarts require retry loops on all statements
zigzag join
phantom reads
triggers
MERGE
grouping sets, cube, rollup
write skew
partial indexes



denormalization
SELECT FOR UPDATE
NULLs in CHECK constraints are truthy
star schemas
transaction contention
sargability
timestamptz doesn't store a timezone
utf8mb4
ascending key problem
ambiguous network errors



cost models don't reflect reality
'null'::jsonb IS NULL = false
TPCC requires wait times
DEFERRABLE INITIALLY IMMEDIATE



cost models don't
reflect reality

EXPLAIN approximates
SELECT COUNT(*)

'null'::jsonb IS NULL = false
DEFERRABLE INITIALLY IMMEDIATE

MATCH PARTIAL foreign keys

TPCC requires wait times

causal reverse



vectorized doesn't
mean SIMD

NULLs are equal in DISTINCT
but inequal in UNIQUE

volcano model

join ordering is NP hard

database cracking

WCOJ

learned indexes

XTID exhaustion



the halloween problem

dee and dum

SERIAL is non-transactional

fsyncgate

allballs

NULL

every sql operator is
actually a join



further learning

- refresher:
<https://www.youtube.com/watch?v=kbKty5ZVKMY>
- pandas experts note:
<https://www.youtube.com/watch?v=fmrmwFPMMaM>
- more discussion:
<https://www.youtube.com/watch?v=OV6Mh2JI9zQ>
- deeper learning:
<https://app.datacamp.com/learn/career-tracks/data-analyst-in-sql>

